

A BEGINNER'S OVERVIEW OF SPARK AND PYSPARK


SAN DIEGO MACHINE LEARNING

MAY 30, 2020



AGENDA

- Provide an overview of what Spark is and why you would use it
- Discuss architecture and key features
- Not a PySpark programming guide
- List some resources to learn more



WHAT IS SPARK

- Apache Spark is an open source general-purpose framework for cluster computing
- History:
 - Originally developed at Berkeley, starting in 2009
 - Contributed to Apache in 2013, and founded Databricks
 - Version 1.0 in 2014. Spark 2.0 in 2016.
- Spark is about computing, not storage or other functionality
- Written in Scala, runs on JVM
- Pros: fast; real-time; powerful caching; language support for Java, Scala, Python, R, and SQL
- Built-in fault tolerance and parallelism


WHY USE SPARK

- Fast, big data computing
- Plays well with HDFS
- Streaming
- Not for multi-user
- MapReduce was inefficient for iterative or interactive jobs
 - According to authors, MapReduce required multiple passes over the data, and multiple batch jobs
- Easy to use




SPARK ARCHITECTURE

- Master node and compute nodes, called driver and executors
- Cluster manager, either built-in, YARN or Mesos
- Everything done in memory – extremely fast
- Functional programming



SPARK ARCHITECTURE

- High level data structures includes DataFrames. DataFrames are made up of partitions.
- Low level data structures are RDDs (Resilient Distributed Datasets)
- Core data structures are immutable, so we use transformations
 - Narrow transformations, within partitions
 - Wide transformations require shuffles




KEY FEATURES

- Low level:
 - RDDs
 - Lazy evaluation, and predicate pushdown
 - Catalyst, cost based optimizer
- High level:
 - DataFrames and Datasets
 - Spark Streaming
- Libraries
 - Spark SQL
 - MLlib
 - Spark NLP

LOW LEVEL FEATURES - RDDS


- RDDs
 - In-memory
 - Fault tolerant
 - Partitioned
 - Immutable
- Can be multiple partitions per worker node
- Often partitioning strategy is matched to data source





LOW LEVEL FEATURES - COMPUTATION

- Computation graph is a tree/DAG
- Transformations
 - Narrow: select, filter, sample, map
 - Wide require “*shuffles*”: aggregate, join, sort
- Lazy evaluation
- Catalyst, cost based optimizer
 - Optimized logical plan, then physical plan
 - Example: predicate pushdown

Narrow Transformations
1 to 1



Wide Transformations (shuffles)
1 to N



- I see many correlations between Spark computation and RDBMS computation

HIGH LEVEL FEATURE - DATAFRAMES


- DataFrames
 - Set of rows with defined schema (not a set of columns)
 - Slow when first introduced, but no longer performance penalty
- Datasets
 - Type-safe data
 - Only available in Scala and Java
- Spark Streaming (separate slide)

LIBRARIES AND PACKAGES

- Spark SQL
 - Large subset of ANSI SQL:2003
- MLlib
 - Similar to NumPy and Scikit-Learn
 - Place numeric features into a vector, then call fit() and transform()
- Spark NLP
- GraphX
- Ecosystem of packages, <https://spark-packages.org>


SPARK STREAMING

- Structured Streaming uses micro-batches
 - Programs for batch and streaming are virtually identical
- Built on top of DataFrames



SPARK STREAMING

- Native support for event time data
- Three output modes:
 - Append
 - Update
 - Complete
- Watermark
 - How late do you expect data to arrive



RESOURCES

- Slides and Python code available at https://github.com/tedkyi/spark_talk
- Spark: The Definitive Guide by Bill Chambers & Matei Zaharia:
<http://shop.oreilly.com/product/0636920034957.do>
- The Apache Spark website documentation:
<https://spark.apache.org/docs/latest/index.html>
- Lots of tutorials and code example on Databricks site:
<https://docs.databricks.com/getting-started/index.html>
- Sign up for your forever free Databricks community account:
<https://community.cloud.databricks.com>

